




learned reward (raw RM)



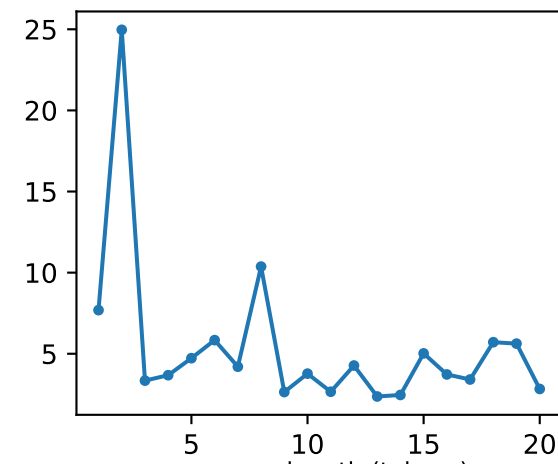
KL to reference (token mean)



policy loss

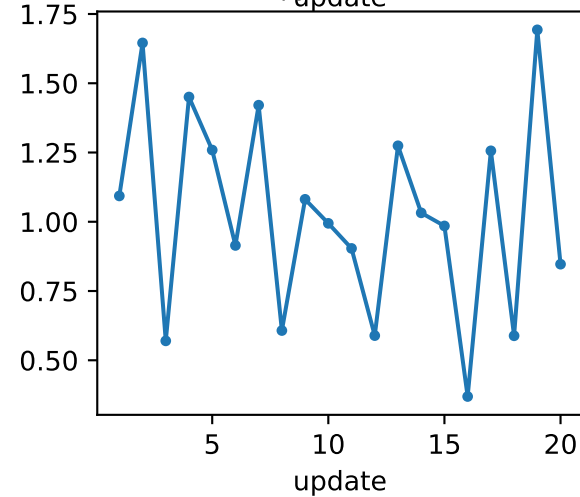


value loss



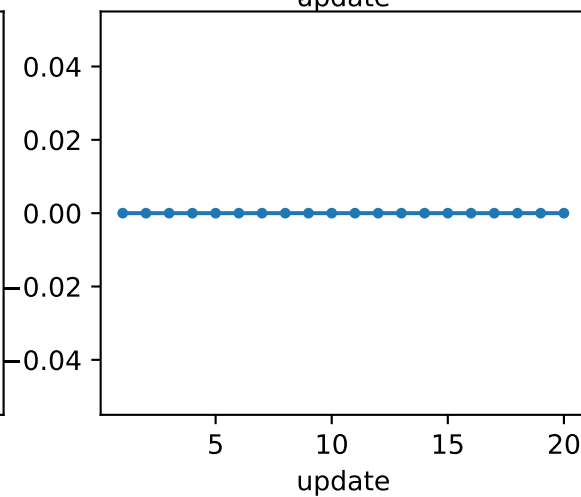
entropy (nats/token)

update



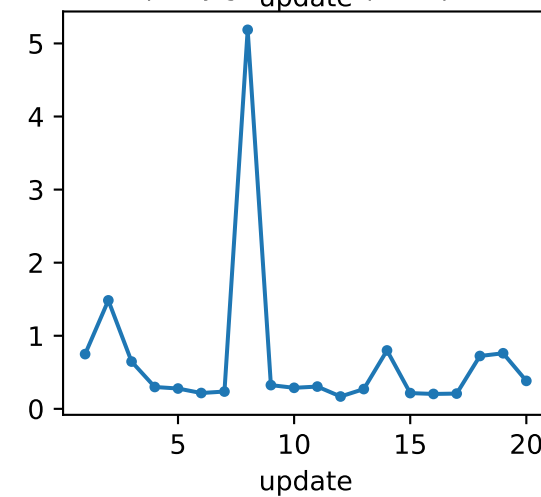
clip fraction

update



policy grad norm (pre-clip)

update



response length (tokens)

update

