




reward (RM)



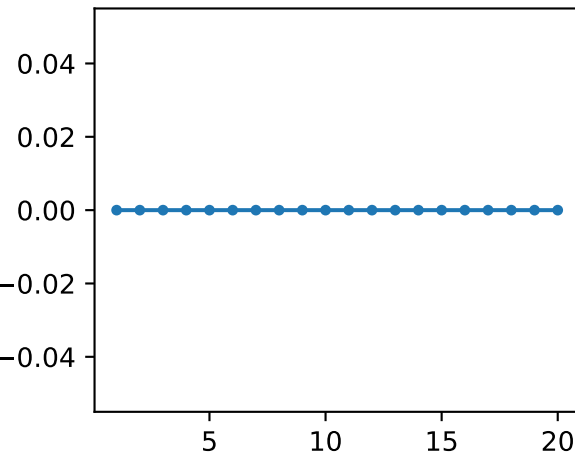
KL to reference (token mean)



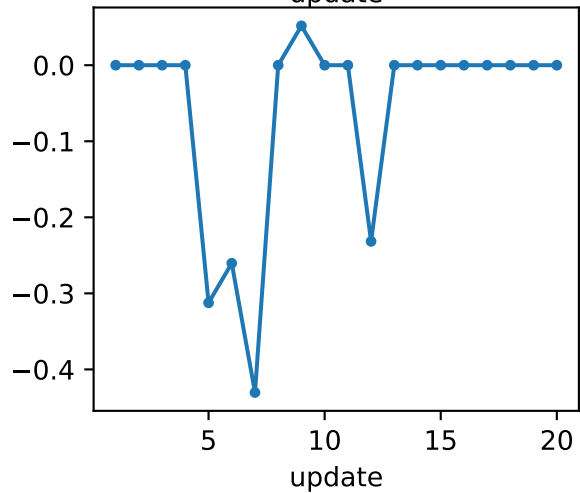
within-group reward std



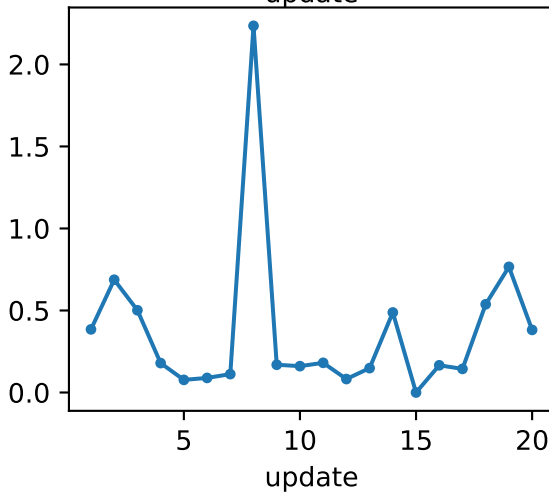
uninformative-group fraction



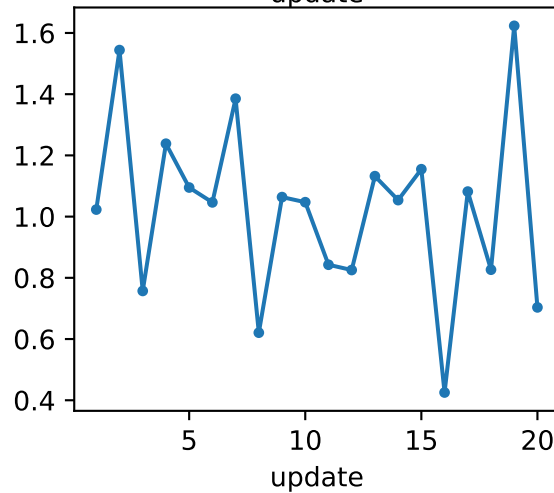
policy loss



grad norm (pre-clip)



entropy (nats/token)



response length (tokens)

